

12.2.1 - Alignment

Artificial Intelligence Policy

Prof. Jack Reilly

S2026

Think:

- *What does it mean for a model to be “aligned”? To whose interests is it aligned, and how do we evaluate that?*

Read:

- **None Required.** Discussion for today will include items originally assigned for [The Singularity and the Future](#) as well; please do the reading originally assigned for that day.

Browse:

- [Claude Constitution](#)
- OpenAI spec: [Introduction](#) and [Source](#)
- [Hidden AI instructions reveal how Anthropic controls Claude 4.](#) 2025. Ars Technica
- Shao et al, 2026. [Future of Work with AI Agents: Auditing Automation and Augmentation Potential across the U.S. Workforce.](#) arXiv.
- Farrell et al, 2025. “[Large AI models are cultural and social technologies](#)”. Science.
- Mobayed, 2025. “[Your Brain on ChatGPT](#)” Psychology Today.
- Bastani et al, 2025. “[Generative AI without guardrails can harm learning: Evidence from high school mathematics](#)” PNAS.
- Berg and Rosenblatt, 2025. “[The Monster inside ChatGPT](#)” Wall Street Journal.
- “[Is AI Rewiring our Minds? Scientists probe cognitive cost of chatbots](#)” Washington Post.
- Li et al, 2025. “(Core Knowledge Deficits in MMLMs)(<https://arxiv.org/abs/2410.10855v4>)” arXiv.
 - *See also:* [Grow AI like a Child](#)
- Gao et al, 2025. [Do Vision-Language Models Have Internal World Models? Towards an Atomic Evaluation](#)

Submit:

- Discussion question to course chat

 Tip

- “ Read”, “ Listen”, and/or “ Watch” items are required content for the day, and should be read/heard/watched before class on that day.
- “ Browse” items should be briefly looked at but do not need to be read deeply unless you want to
- “ Additional Resources” do not need to be looked at; they are there to serve, if useful, as further references for your debates, final projects, and general edification later.