

12.2 - Alignment

Artificial Intelligence Policy

Prof. Jack Reilly

S2026

In Progress

Think:

Read:

Browse:

- claude constitution
- openai spec
- [Hidden AI instructions reveal how Anthropic controls Claude 4](#). 2025. Ars Technica
- Shao et al, 2026. [Future of Work with AI Agents: Auditing Automation and Augmentation Potential across the U.S. Workforce](#). arXiv.
- Farrell et al, 2025. [“Large AI models are cultural and social technologies”](#). Science.
- Mobayed, 2025. [“Your Brain on ChatGPT”](#) Psychology Today.
- Bastani et al, 2025. [“Generative AI without guardrails can harm learning: Evidence from high school mathematics”](#) PNAS.
- Berg and Rosenblatt, 2025. [“The Monster inside ChatGPT”](#) Wall Street Journal.
- [“Is AI Rewiring our Minds? Scientists probe cognitive cost of chatbots”](#) Washington Post.
- Li et al, 2025. [“\(Core Knowledge Deficits in MMLMs\)\(https://arxiv.org/abs/2410.10855v4\)”](#) arXiv.
 - *See also:* [Grow AI like a Child](#)
- Gao et al, 2025. [Do Vision-Language Models Have Internal World Models? Towards an Atomic Evaluation](#)

Submit:

- Discussion question to course chat

 Tip

- “ Read”, “ Listen”, and/or “ Watch” items are required content for the day, and should be read/heard/watched before class on that day.
- “ Browse” items should be briefly looked at but do not need to be read deeply unless you want to
- “ Additional Resources” do not need to be looked at; they are there to serve, if useful, as further references for your debates, final projects, and general edification later.